

GMetis - Xeon Phi

David Pereira

pg22821@alunos.uminho.pt

Rui Brito

pg22781@alunos.uminho.pt

Austin, August 2013

Abstract

1 Introduction

- GMetis is a graph partitioning application which uses the Galois framework
- Consists of three major phases
 - Coarsening
 - * Find matching nodes
 - * Create Coarse Edges
 - Initial Partitioning (Clustering)
 - Refinement
- Given a graph $G_0 = (V_0, E_0)$:
 - Coarsening
 - * G_0 is transformed into a sequence of smaller graphs G_1, G_2, \dots, G_m such that $|V_0| > |V_1| > |V_2| > \dots > |V_m|$
 - Partitioning
 - * A 2-way partition P_m of the graph $G_m = (V_m, E_m)$ is computed that partitions V_m into two parts, each containing half the vertices of G_0
 - Refinement
 - * The partition P_m of G_m is projected back to G_0 by going through intermediate partitions $P_{m-1}, P_{m-2}, \dots, P_1, P_0$

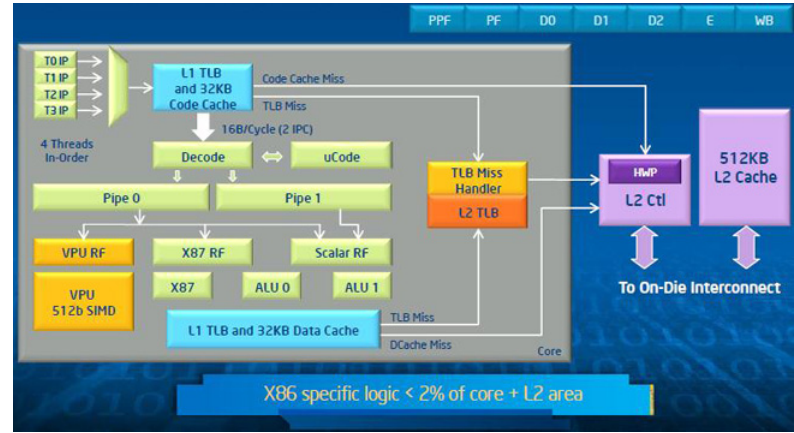
2 System characteristics

The measurements were performed in both Stampede's hosts and coprocessors. The hosts are comprised of dual Intel Xeon E5-2680, while the coprocessors are the new Intel Xeon Phi with 61 cores. Their characteristics are presented in the following tables.

Apart from the characteristics showed in table ??, there are others that should be mentioned. Each core contains a in-order dual pipeline which can issue two instructions from the same hardware thread per clock cycle. However, the front-end of the pipeline does not issue instructions from the same hardware thread in consecutive cycles. This means that the maximum issue rate is only attainable with at least 2 threads per core

while the other threads have the purpose of hiding pipeline stalls due to memory latency.

The fact that the pipeline issue instructions in-order increases memory related problems.



3 Metis

4 Mt-metis




Figure 1: Mt-metis - 128 partitions

5 GMetis and the Galois Framework




Figure 2: GMetis - 128 partitions

This figure shows the scalability of gmetis on Xeon phi over the runtime with on thread. This example is for 128 partition, but we did measurements for different number of partitons, such as 16 and 1024. All of them have a similar behaviour.

6 Conclusion

Results showed that both *Metis* and *Mt-metis* have better edgecut than *Gmetis*. However, Gmetis’s runtime is lower for a high number of partitions.

Xeon Phi provides a theoretical performance of 2112 GFlop/sec for double precision arithmetic and 1056 GFlop-s/sec for single precision arithmetic. This values comprises the use of 60 cores since one core is necessary to perform operating system operations.

References

Manufacturer	Intel
Model	Xeon E5-2680
μ Arch	Sandy Bridge
Clock freq	2.70 GHz
#CPUs (sockets)	2
#Cores/CPU	8
#Thread/Core	1
L1 cache size/core	32 KB
L2 cache size/core	256 KB
L3 shared cache size/CPU	20 MB
Main Memory/CPU	16 GB
Vector width	256 bits (AVX)

Table 1: Intel Xeon E5-2680

Manufacturer	Intel
Model	Xeon Phi SE10P
μ Arch	Many Integrated Cores - MIC
Clock freq	1.1 GHz
#CPUs (sockets)	1
#Cores/CPU	61
#Thread/Core	4
L1 cache size/core	32KB
L2 cache size/core	512 KB
Main Memory/CPU	8 GB
Vector width	512 bits

Table 2: Intel Xeon Phi